

Types of Indexes and partitions used in Data warehousing

Indexes

- 1) B-tree Indexes: B-tree indexes are the most common type of index used in data warehousing. They organize data in a balanced tree structure, allowing for efficient retrieval of data based on key values. B-tree indexes are well-suited for range queries and equality searches. They are particularly useful in OLAP (Online Analytical Processing) environments where ad-hoc queries are common.
- Reverse Key Indexes:
Reverse key indexes are a type of B-tree index where the bytes of the index key are reversed before being stored in the index. They are particularly useful in environments where index key values are monotonically increasing, such as in timestamp or sequence-based columns.
Reverse key indexes help distribute insert activity across the index structure, reducing contention and hotspots, which can improve performance in data warehousing scenarios with high insert rates.

2)Bitmap Indexes: Bitmap indexes are highly efficient for columns with low cardinality, meaning columns that have a relatively small number of distinct values. Instead of storing pointers to rows like B-tree indexes, bitmap indexes store a bitmap for each distinct value in the column. Each bit in the bitmap corresponds to a row in the table, indicating whether the row contains the specific value or not. Bitmap indexes are beneficial for data warehousing applications with queries involving multiple columns, such as data mining and decision support systems.

3)Function-Based Indexes: Function-based indexes are created based on expressions or functions applied to columns in a table. These indexes allow for efficient retrieval of data when queries involve transformations or computations on columns. Function-based indexes are particularly useful in data warehousing environments where data is pre-aggregated or transformed before analysis.

4)Bitmap Join Indexes:

Bitmap join indexes are a specialized type of bitmap index used to improve the performance of join operations in data warehousing environments.

They precompute join operations between two or more tables and store the result as bitmaps, representing the presence or absence of matching rows.

Bitmap join indexes are particularly effective for star and snowflake schema designs commonly used in data warehousing, where join queries involving large fact tables and dimension tables are prevalent.

Partitions

- ✚ Range Partitions: Range partitions divide data based on a specified range of values. For example, data can be partitioned based on date ranges, where each partition contains data for a specific time period (e.g., monthly or quarterly partitions). Range partitions are beneficial for data warehousing applications with time-series data, as they facilitate efficient data pruning and retention policies.

✚ Hash Partitions: Hash partitions distribute data evenly across multiple partitions based on a hash function applied to a specified column or set of columns. Hash partitions ensure data is evenly distributed, which helps balance query workload across partitions. They are commonly used in data warehousing environments with large, evenly distributed datasets.

✚ List Partitions: List partitions allow data to be grouped into partitions based on a predefined list of values for a specific column. Each partition contains rows with values matching the specified list. List partitions are useful for categorizing data into discrete groups, such as geographical regions or product categories.

✚ Composite Partitions: Composite partitions combine multiple partitioning methods (e.g., range and hash) to create a hierarchical partitioning scheme. For example, a table may be range-partitioned by date and then hash-partitioned within each range partition. Composite partitions offer flexibility in organizing data based on multiple criteria, optimizing both data retrieval and management.